

Code AXA

October 23, 2024

```
[1]: # Importing libraries

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: # Load the dataset

file_path = 'depression_data.csv'
depression_data = pd.read_csv(file_path)
```

```
[3]: # Display the first few rows of the dataset to understand its structure
depression_data.head()
```

```
[3]:
```

	Name	Age	Marital Status	Education Level	\
0	Christine Barker	31	Married	Bachelor's Degree	
1	Jacqueline Lewis	55	Married	High School	
2	Shannon Church	78	Widowed	Master's Degree	
3	Charles Jordan	58	Divorced	Master's Degree	
4	Michael Rich	18	Single	High School	

	Number of Children	Smoking Status	Physical Activity Level	\
0	2	Non-smoker	Active	
1	1	Non-smoker	Sedentary	
2	1	Non-smoker	Sedentary	
3	3	Non-smoker	Moderate	
4	0	Non-smoker	Sedentary	

	Employment Status	Income	Alcohol Consumption	Dietary Habits	\
0	Unemployed	26265.67	Moderate	Moderate	
1	Employed	42710.36	High	Unhealthy	
2	Employed	125332.79	Low	Unhealthy	
3	Unemployed	9992.78	Moderate	Moderate	
4	Unemployed	8595.08	Low	Moderate	

	Sleep Patterns	History of Mental Illness	History of Substance Abuse	\
--	----------------	---------------------------	----------------------------	---

0	Fair	Yes	No
1	Fair	Yes	No
2	Good	No	No
3	Poor	No	No
4	Fair	Yes	No

Family History of Depression Chronic Medical Conditions			
0		Yes	Yes
1		No	Yes
2		Yes	No
3		No	No
4		Yes	Yes

```
[4]: # Basic statistics and missing values check
summary_stats = depression_data.describe(include='all')
print(summary_stats)
```

	Name	Age	Marital Status	Education Level \
count	413768	413768.000000	413768	413768
unique	196851	NaN	4	5
top	Michael Smith	NaN	Married	Bachelor's Degree
freq	198	NaN	240444	124329
mean	NaN	49.000713	NaN	NaN
std	NaN	18.158759	NaN	NaN
min	NaN	18.000000	NaN	NaN
25%	NaN	33.000000	NaN	NaN
50%	NaN	49.000000	NaN	NaN
75%	NaN	65.000000	NaN	NaN
max	NaN	80.000000	NaN	NaN

	Number of Children	Smoking Status	Physical Activity Level \
count	413768.000000	413768	413768
unique	NaN	3	3
top	NaN	Non-smoker	Sedentary
freq	NaN	247416	176850
mean	1.298972	NaN	NaN
std	1.237054	NaN	NaN
min	0.000000	NaN	NaN
25%	0.000000	NaN	NaN
50%	1.000000	NaN	NaN
75%	2.000000	NaN	NaN
max	4.000000	NaN	NaN

	Employment Status	Income	Alcohol Consumption	Dietary Habits \
count	413768	413768.000000	413768	413768
unique	2	NaN	3	3
top	Employed	NaN	Moderate	Unhealthy
freq	265659	NaN	173440	170817

mean	NaN	50661.707971	NaN	NaN
std	NaN	40624.100565	NaN	NaN
min	NaN	0.410000	NaN	NaN
25%	NaN	21001.030000	NaN	NaN
50%	NaN	37520.135000	NaN	NaN
75%	NaN	76616.300000	NaN	NaN
max	NaN	209995.220000	NaN	NaN

	Sleep Patterns	History of Mental Illness	History of Substance Abuse	\
count	413768	413768	413768	
unique	3	2	2	
top	Fair	No	No	
freq	196789	287943	284880	
mean	NaN	NaN	NaN	
std	NaN	NaN	NaN	
min	NaN	NaN	NaN	
25%	NaN	NaN	NaN	
50%	NaN	NaN	NaN	
75%	NaN	NaN	NaN	
max	NaN	NaN	NaN	

	Family History of Depression	Chronic Medical Conditions
count	413768	413768
unique	2	2
top	No	No
freq	302515	277561
mean	NaN	NaN
std	NaN	NaN
min	NaN	NaN
25%	NaN	NaN
50%	NaN	NaN
75%	NaN	NaN
max	NaN	NaN

```
[5]: # Check for missing values
missing_values = depression_data.isnull().sum()
print(missing_values)
```

Name	0
Age	0
Marital Status	0
Education Level	0
Number of Children	0
Smoking Status	0
Physical Activity Level	0
Employment Status	0
Income	0
Alcohol Consumption	0

```

Dietary Habits          0
Sleep Patterns          0
History of Mental Illness 0
History of Substance Abuse 0
Family History of Depression 0
Chronic Medical Conditions 0
dtype: int64

```

```

[6]: # Set plot style
sns.set(style="whitegrid")

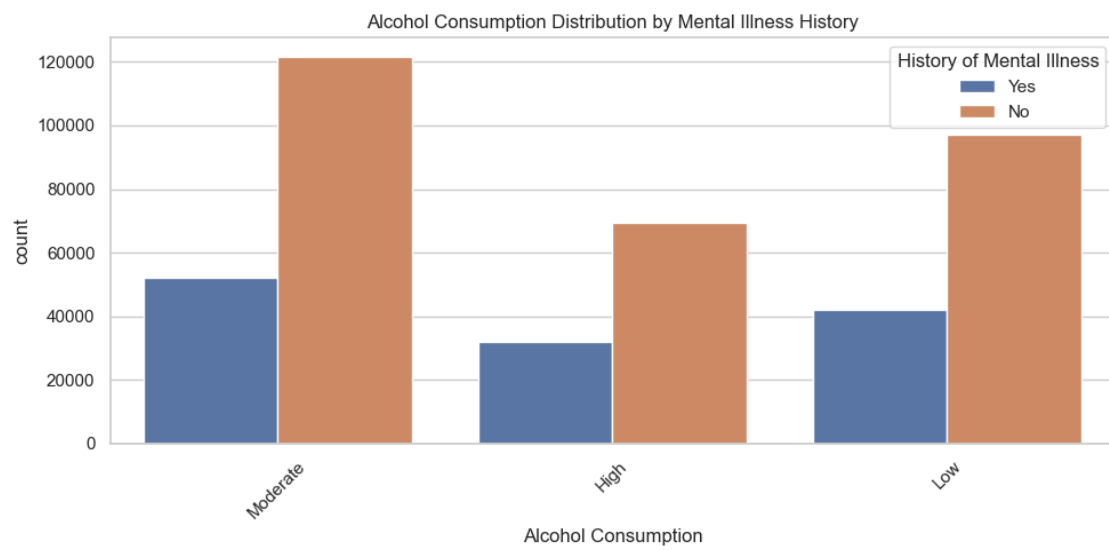
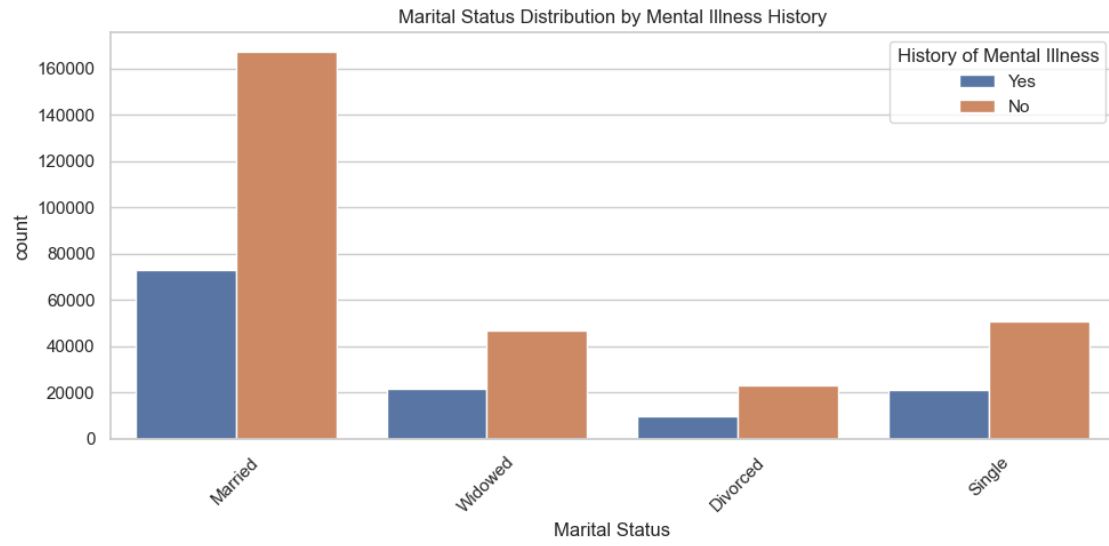
# Distribution of Marital Status
plt.figure(figsize=(10, 5))
sns.countplot(x='Marital Status', hue='History of Mental Illness',
↳data=depression_data)
plt.title('Marital Status Distribution by Mental Illness History')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()


# Alcohol consumption comparison based on mental illness history
plt.figure(figsize=(10, 5))
sns.countplot(x='Alcohol Consumption', hue='History of Mental Illness',
↳data=depression_data)
plt.title('Alcohol Consumption Distribution by Mental Illness History')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Correlation heatmap for numerical data
numerical_data = depression_data[['Age', 'Income', 'Number of Children']]
correlation = numerical_data.corr()

plt.figure(figsize=(8, 6))
sns.heatmap(correlation, annot=True, cmap="coolwarm", linewidths=0.5)
plt.title('Correlation Heatmap')
plt.tight_layout()
plt.show()

```





Key Insights from Exploratory Data Analysis: Marital Status:

Individuals who are married seem to have a slightly lower prevalence of mental illness history compared to those who are widowed, divorced, or single. Alcohol Consumption:

Moderate alcohol consumption appears more common among individuals without a history of mental illness, whereas individuals with a history of mental illness are more spread out across low and moderate consumption categories. Correlation Heatmap:

The correlation between numerical variables such as age, income, and number of children is weak, suggesting these features may not strongly interact with each other. However, these may still be useful for predicting mental illness when combined with other factors.

```
[7]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

# Convert categorical variables to numerical using Label Encoding
label_encoders = {}
for column in depression_data.select_dtypes(include=['object']).columns:
    label_encoders[column] = LabelEncoder()
```

```

    depression_data[column] = label_encoders[column].
    ↪fit_transform(depression_data[column])

# Split the data into features and target variable
X = depression_data.drop('History of Mental Illness', axis=1)
y = depression_data['History of Mental Illness']

# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ↪
    ↪random_state=42)

# We'll start by trying a Random Forest Classifier and a Logistic Regression ↪
    ↪model to compare their performance
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

# Random Forest Classifier
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)

# Logistic Regression
lr_model = LogisticRegression(max_iter=1000, random_state=42)
lr_model.fit(X_train, y_train)
lr_predictions = lr_model.predict(X_test)

# Performance report for both models
rf_report = classification_report(y_test, rf_predictions)
lr_report = classification_report(y_test, lr_predictions)

rf_report, lr_report

```

```

[7]: (
      precision    recall  f1-score   support\n\n          0
0.70      0.95      0.80   57471\n          1      0.38      0.07      0.12
25283\n\n      accuracy                0.68   82754\n      macro avg
0.54      0.51      0.46   82754\nweighted avg          0.60      0.68      0.60
82754\n',
      '
      precision    recall  f1-score   support\n\n          0
0.69      1.00      0.82   57471\n          1      0.00      0.00      0.00
25283\n\n      accuracy                0.69   82754\n      macro avg
0.35      0.50      0.41   82754\nweighted avg          0.48      0.69      0.57
82754\n')

```

```

[8]: from imblearn.over_sampling import SMOTE

```

```
[9]: # Apply SMOTE to the training data to handle class imbalance
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

[10]: # Train the Random Forest model again on the balanced dataset
random_forest_smote = RandomForestClassifier(n_estimators=100, random_state=42)
random_forest_smote.fit(X_train_smote, y_train_smote)

[10]: RandomForestClassifier(random_state=42)

[11]: # Make predictions on the test set
rf_preds_smote = random_forest_smote.predict(X_test)

[12]: # Evaluate the Random Forest model on the test data after SMOTE
from sklearn.metrics import classification_report, roc_auc_score
rf_report_smote = classification_report(y_test, rf_preds_smote)
rf_auc_smote = roc_auc_score(y_test, rf_preds_smote)

rf_report_smote, rf_auc_smote

[12]: ( '          precision    recall  f1-score   support\n\n
0.71      0.78      0.74   57471\n          1      0.34      0.27      0.30
25283\n\n accuracy              0.62   82754\n macro avg
0.53      0.52      0.52   82754\nweighted avg          0.60      0.62      0.61
82754\n',
0.5215343051290768)
```

0.1 Hyperparameter Tuning

```
[13]: from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

# Define the parameter distribution
param_dist = {
    'n_estimators': randint(100, 500),
    'max_depth': [10, 20, 30, None],
    'min_samples_split': randint(2, 10),
    'min_samples_leaf': randint(1, 5)
}

# Initialize the Random Forest model
random_forest = RandomForestClassifier(random_state=42)

# Set up RandomizedSearchCV
random_search = RandomizedSearchCV(estimator=random_forest,
    param_distributions=param_dist,
```



```

n_iter=10, cv=2, n_jobs=-1, verbose=2,
random_state=42, scoring='roc_auc')

# Fit the random search to the data (using SMOTE-balanced data)
random_search.fit(X_train_smote, y_train_smote)

# Get the best parameters from the search
best_params = random_search.best_params_
print(f"Best parameters found: {best_params}")

# Evaluate the tuned model on the test set
best_rf_model = random_search.best_estimator_
rf_preds_tuned = best_rf_model.predict(X_test)

# Get the evaluation report
rf_report_tuned = classification_report(y_test, rf_preds_tuned)
rf_auc_tuned = roc_auc_score(y_test, rf_preds_tuned)

print(rf_report_tuned)
print(f"ROC-AUC score after tuning: {rf_auc_tuned}")

```

```

Fitting 2 folds for each of 10 candidates, totalling 20 fits
[CV] END max_depth=30, min_samples_leaf=3, min_samples_split=9,
n_estimators=472; total time= 2.3min
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2,
n_estimators=158; total time= 47.7s
[CV] END max_depth=30, min_samples_leaf=4, min_samples_split=6,
n_estimators=370; total time= 1.8min
[CV] END max_depth=None, min_samples_leaf=4, min_samples_split=6,
n_estimators=260; total time= 1.3min
[CV] END max_depth=30, min_samples_leaf=3, min_samples_split=9,
n_estimators=472; total time= 2.5min
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2,
n_estimators=158; total time= 49.6s
[CV] END max_depth=None, min_samples_leaf=4, min_samples_split=9,
n_estimators=230; total time= 1.1min
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=7,
n_estimators=485; total time= 2.3min
[CV] END max_depth=30, min_samples_leaf=2, min_samples_split=4,
n_estimators=314; total time= 1.8min
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=7,
n_estimators=352; total time= 1.6min
[CV] END max_depth=None, min_samples_leaf=4, min_samples_split=9,
n_estimators=230; total time= 1.2min
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=7,
n_estimators=485; total time= 2.4min
[CV] END max_depth=30, min_samples_leaf=4, min_samples_split=6,
n_estimators=370; total time= 2.0min

```

[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=7,
n_estimators=352; total time= 1.7min

Best parameters found: {'max_depth': 20, 'min_samples_leaf': 1,
'min_samples_split': 3, 'n_estimators': 443}

	precision	recall	f1-score	support
0	0.71	0.79	0.74	57471
1	0.35	0.25	0.29	25283
accuracy			0.63	82754
macro avg	0.53	0.52	0.52	82754
weighted avg	0.60	0.63	0.61	82754

ROC-AUC score after tuning: 0.5209567051260741

1 Ensemble Techniques

1.1 XGBoost

```
[14]: import xgboost as xgb
from sklearn.metrics import classification_report, roc_auc_score
from xgboost import XGBClassifier

# Initialize the XGBoost model
xgb_model = XGBClassifier(
    objective='binary:logistic',
    scale_pos_weight=len(y_train_smote[y_train_smote == 0]) /
    len(y_train_smote[y_train_smote == 1]), # Handling class imbalance
    use_label_encoder=False,
    eval_metric='logloss', # A common evaluation metric for binary
    classification
    random_state=42
)

# Fit the model
xgb_model.fit(X_train_smote, y_train_smote)

# Make predictions on the test data
xgb_preds = xgb_model.predict(X_test)

# Evaluate the model
xgb_report = classification_report(y_test, xgb_preds)
xgb_auc = roc_auc_score(y_test, xgb_preds)

# Display the report and AUC
print("XGBoost Classification Report:")
print(xgb_report)
```

```
print(f"ROC-AUC Score: {xgb_auc}")
```

XGBoost Classification Report:

	precision	recall	f1-score	support
0	0.70	0.82	0.76	57471
1	0.35	0.22	0.27	25283
accuracy			0.64	82754
macro avg	0.53	0.52	0.51	82754
weighted avg	0.60	0.64	0.61	82754

ROC-AUC Score: 0.5198005567630579

1.2 LightGBM

```
[15]: import lightgbm as lgb
from sklearn.metrics import classification_report, roc_auc_score
from lightgbm import LGBMClassifier

# Initialize the LightGBM model
lgb_model = LGBMClassifier(
    objective='binary',
    is_unbalance=True, # Automatically balances classes
    random_state=42,
    n_estimators=1000, # Number of boosting rounds
    learning_rate=0.05, # Learning rate to control the size of a step during
    ↪ optimization
    max_depth=20, # Maximum depth of the trees
)

# Train the model
lgb_model.fit(X_train_smote, y_train_smote)

# Make predictions on the test set
lgb_preds = lgb_model.predict(X_test)

# Evaluate the model
lgb_report = classification_report(y_test, lgb_preds)
lgb_auc = roc_auc_score(y_test, lgb_preds)

# Print the classification report and ROC-AUC score
print("LightGBM Classification Report:")
print(lgb_report)
print(f"ROC-AUC Score: {lgb_auc}")
```

[LightGBM] [Warning] Accuracy may be bad since you didn't explicitly set num_leaves OR $2^{\text{max_depth}} > \text{num_leaves}$. (num_leaves=31).

```
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Warning] Accuracy may be bad since you didn't explicitly set
num_leaves OR 2^max_depth > num_leaves. (num_leaves=31).
[LightGBM] [Info] Number of positive: 230472, number of negative: 230472
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.007958 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] Number of data points in the train set: 460944, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Warning] Accuracy may be bad since you didn't explicitly set
num_leaves OR 2^max_depth > num_leaves. (num_leaves=31).
LightGBM Classification Report:
```

	precision	recall	f1-score	support
0	0.70	0.84	0.77	57471
1	0.35	0.19	0.25	25283
accuracy			0.64	82754
macro avg	0.53	0.52	0.51	82754
weighted avg	0.60	0.64	0.61	82754

ROC-AUC Score: 0.5182308621161287

1.3 Stacking Model

```
[16]: from sklearn.ensemble import StackingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
from lightgbm import LGBMClassifier
from sklearn.metrics import classification_report, roc_auc_score

# Define the base models
base_models = [
    ('random_forest', RandomForestClassifier(n_estimators=200,
    random_state=42)),
    ('xgboost', xgb.XGBClassifier(use_label_encoder=False,
    eval_metric='logloss', random_state=42)),
    ('lightgbm', LGBMClassifier(random_state=42))
]

# Define the meta-model
meta_model = LogisticRegression()
```

```

# Create the Stacking Classifier
stacking_model = StackingClassifier(estimators=base_models,
    ↪final_estimator=meta_model, cv=5, n_jobs=-1)

# Train the Stacking model using the SMOTE balanced data
stacking_model.fit(X_train_smote, y_train_smote)

# Make predictions on the test set
stacking_preds = stacking_model.predict(X_test)

# Evaluate the Stacking model
stacking_report = classification_report(y_test, stacking_preds)
stacking_auc = roc_auc_score(y_test, stacking_preds)

# Display the report and ROC-AUC score
print("Stacking Classifier Report:")
print(stacking_report)
print(f"ROC-AUC Score: {stacking_auc}")

```

```

[CV] END max_depth=30, min_samples_leaf=2, min_samples_split=4,
n_estimators=314; total time= 1.6min
[CV] END max_depth=None, min_samples_leaf=4, min_samples_split=6,
n_estimators=260; total time= 1.3min
[CV] END max_depth=30, min_samples_leaf=4, min_samples_split=6,
n_estimators=120; total time= 38.7s
[CV] END max_depth=20, min_samples_leaf=1, min_samples_split=3,
n_estimators=443; total time= 2.4min
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Info] Number of positive: 184378, number of negative: 184377
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Info] Number of positive: 184378, number of negative: 184378
[LightGBM] [Info] Number of positive: 184378, number of negative: 184377
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Info] Number of positive: 184377, number of negative: 184378
[LightGBM] [Info] Number of positive: 184377, number of negative: 184378
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.105127 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] Number of data points in the train set: 368755, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500001 -> initscore=0.000005
[LightGBM] [Info] Start training from score 0.000005
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.152572 seconds.

```

You can set `force_col_wise=true` to remove the overhead.

```
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] Number of data points in the train set: 368756, number of used
features: 15
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.124604 seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
[LightGBM] [Info] Number of data points in the train set: 368755, number of used
features: 15
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.171513 seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] Number of data points in the train set: 368755, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500001 -> initscore=0.000005
[LightGBM] [Info] Start training from score 0.000005
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499999 -> initscore=-0.000005
[LightGBM] [Info] Start training from score -0.000005
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.184576 seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] Number of data points in the train set: 368755, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499999 -> initscore=-0.000005
[LightGBM] [Info] Start training from score -0.000005
```

Stacking Classifier Report:

	precision	recall	f1-score	support
0	0.71	0.79	0.75	57471
1	0.35	0.26	0.30	25283
accuracy			0.63	82754
macro avg	0.53	0.52	0.52	82754
weighted avg	0.60	0.63	0.61	82754

ROC-AUC Score: 0.5232308511288138

1.4 Adjusting Class Weights in the Stacking Classifier

```
[17]: from sklearn.utils.class_weight import compute_class_weight
import numpy as np

# Calculate class weights
class_weights = compute_class_weight(class_weight='balanced', classes=np.
    unique(y_train_smote), y=y_train_smote)
class_weights_dict = {0: class_weights[0], 1: class_weights[1]}

# Update the base models with class weights
base_models_weighted = [
    ('random_forest', RandomForestClassifier(n_estimators=200, random_state=42,
    class_weight=class_weights_dict)),
    ('xgboost', xgb.XGBClassifier(use_label_encoder=False,
    eval_metric='logloss', random_state=42,
    scale_pos_weight=class_weights_dict[1])),
    ('lightgbm', LGBMClassifier(random_state=42,
    class_weight=class_weights_dict))
]

# Create the Stacking Classifier with class weights
stacking_model_weighted = StackingClassifier(estimators=base_models_weighted,
    final_estimator=LogisticRegression(class_weight=class_weights_dict), cv=5,
    n_jobs=-1)

# Train the model
stacking_model_weighted.fit(X_train_smote, y_train_smote)

# Make predictions on the test set
stacking_preds_weighted = stacking_model_weighted.predict(X_test)

# Evaluate the model
stacking_report_weighted = classification_report(y_test,
    stacking_preds_weighted)
stacking_auc_weighted = roc_auc_score(y_test, stacking_preds_weighted)

# Display the updated report and ROC-AUC score
print("Stacking Classifier with Class Weights Report:")
print(stacking_report_weighted)
print(f"ROC-AUC Score: {stacking_auc_weighted}")
```

```
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Info] Number of positive: 184378, number of negative: 184377
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Info] Number of positive: 184377, number of negative: 184378
```

```

[LightGBM] [Info] Number of positive: 184377, number of negative: 184378
[LightGBM] [Info] Number of positive: 184378, number of negative: 184377
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Info] Number of positive: 184378, number of negative: 184378
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.058180 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] Number of data points in the train set: 368755, number of used
features: 15
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.085135 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] Number of data points in the train set: 368755, number of used
features: 15
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.063318 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500001 -> initscore=0.000005
[LightGBM] [Info] Start training from score 0.000005
[LightGBM] [Info] Number of data points in the train set: 368755, number of used
features: 15
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.072224 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499999 -> initscore=-0.000005
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500001 -> initscore=0.000005
[LightGBM] [Info] Start training from score -0.000005
[LightGBM] [Info] Number of data points in the train set: 368755, number of used
features: 15
[LightGBM] [Info] Start training from score 0.000005
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499999 -> initscore=-0.000005
[LightGBM] [Info] Start training from score -0.000005
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.086936 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 611
[LightGBM] [Info] Number of data points in the train set: 368756, number of used
features: 15
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000
Stacking Classifier with Class Weights Report:
      precision    recall  f1-score   support

```


0	0.71	0.79	0.75	57471
1	0.35	0.26	0.30	25283
accuracy			0.63	82754
macro avg	0.53	0.52	0.52	82754
weighted avg	0.60	0.63	0.61	82754

ROC-AUC Score: 0.5232308511288138

[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines

[LightGBM] [Info] Number of positive: 230472, number of negative: 230472

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.014339 seconds.

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 611

[LightGBM] [Info] Number of data points in the train set: 460944, number of used features: 15

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[CV] END max_depth=30, min_samples_leaf=4, min_samples_split=6,

n_estimators=120; total time= 35.4s

[CV] END max_depth=20, min_samples_leaf=1, min_samples_split=3,

n_estimators=443; total time= 2.2min

[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines

[LightGBM] [Info] Number of positive: 230472, number of negative: 230472

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.023533 seconds.

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 611

[LightGBM] [Info] Number of data points in the train set: 460944, number of used features: 15

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000

[]: